


P2SD: shared image computation, then compact per-prompt prediction



Blue: trained modules Green: codes / context Gray: data or frozen (explicitly labeled) Gold: losses

No full-resolution sheet decode for the latent training loss. The binary branch still decodes one crop per CT batch.