

PLatS: Papyrus Surface Segmentation in Latent Space

Houjing Huang

houjing.huang@gmail.com

Abstract

Separating tightly packed papyrus sheets requires preserving their geometry and identity: a small false-positive bridge can join neighboring layers even when voxel overlap is high. We present PLatS, which learns a single-sheet latent space and predicts its codes from CT and positive point prompts. A transformer trained with latent MSE uses the same representation for sheet reconstruction and automatic seed clustering, postponing dense decoding until instances are selected. Across 106 released Kaggle test cubes, 87.3% of 764 eight-point sheet predictions have one component and no tunnels or cavities before cleanup; single-point codes distinguish annotated sheet identities with 0.912 AUROC. Automatic union scores remain below the Kaggle winner. We analyze topology errors caused by ignore-mask erasure and demonstrate independently flattened, centimeter-scale CT and ink outputs. The results support a useful sheet prior while exposing incomplete coverage, instance errors and stitching drift.

1. Introduction

Virtual unwrapping needs surfaces that stay on one papyrus sheet. In densely packed CT volumes, a few erroneous voxels can connect neighboring layers; removing those voxels after binary segmentation requires recovering which sheet they belong to. Touching sheets remain a difficulty in the Kaggle winner’s account [2]. This motivates predicting individual sheet geometry together with a representation of sheet identity.

Our central idea is to use *one latent code for reconstruction and instance discovery*. A denoising autoencoder learns a compact representation of individual sheets. Given CT and one or more points, a transformer regresses that representation with MSE. The decoded code reconstructs a sheet; distances between codes group automatically sampled foreground points. Prompt-dependent computation stays on a compact latent grid; dense sheet decoding is required only for retained instance hypotheses.

We contribute (i) a point-conditioned sheet representa-

tion that couples a learned shape prior with latent identity supervision; (ii) an automatic seed–code–cluster–decode procedure that shares CT computation; and (iii) an evaluation separating raw sheet topology, automatic instance recovery, and topology changes introduced by annotation-ignorance erasure. Large independently flattened surfaces expose downstream coverage and drift. The representation empirically favors valid sheet topology; it provides no mathematical guarantee. Automatic segmentation does not yet exceed the winner’s combined score.

Relation to prior work. ACNN learns anatomical representations to regularize image analysis [3]; GSPN reconstructs shape proposals for point-cloud instances [8]. PointSAM studies promptable 3D point-cloud segmentation [9]. PLatS applies a shared reconstruction and identity code to thin volumetric sheets, with clustering before dense decoding. Betti matching provides spatially matched topology losses and metrics [6]; here it is used for evaluation, while the sheet prior is learned without an explicit topology loss. These connections motivate the design; they are not matched baseline comparisons.

2. Method

PLatS learns a sheet representation, then learns to predict it from CT and points (Figs. 1–2). At inference, the same representation groups foreground seeds and reconstructs individual sheets.

2.1. Stage 1: learn the sheet AE

Let S be a clean binary sheet mask and $\tilde{S}$ a corrupted version. The encoder E maps it to a latent code z , and the occupancy decoder D reconstructs a probability map Q :

$$z = E(\tilde{S}), \quad Q = \text{sigmoid}(D(z)). \quad (1)$$

Here sigmoid is the sigmoid function and D includes the final occupancy head. There are no encoder–decoder skip connections: reconstruction must pass through z . Intermediate occupancy heads predict $Q^{(a)}$, where a indexes decoder stages. A sibling dense head predicts unsigned sheet




Figure 1. **Learning the sheet representation.** Actual corrupted input, reconstruction and clean target use matched slices. Occupancy and distance predictions connect to their losses and clean targets; the query head uses the same distance target at coordinate q . The prominent code is the only path through the AE. Below, two real sheets illustrate cosine-margin code repulsion (measured cosine 0.332). Masks are displayed at threshold 0.5.

distance U ; a coordinate-query head $h(z, q)$ predicts distance at location q .

Final and intermediate occupancy heads use binary cross-entropy (BCE) and soft Dice losses against the clean sheet. Dense and query distances use Smooth L1 losses against GT distance targets. A cosine-margin penalty repels codes of different sheets; spatial penalties suppress foreground outside the sheet and at the crop border. All AE modules train in this stage and are then frozen. The bottleneck *empirically favors valid sheet topology*, without a mathematical guarantee or an explicit topology loss.

2.2. Stage 2: learn point-to-latent prediction

For CT crop X and positive point set P , the image encoder and context blocks C produce shared features F . The point transformer and code heads T predict a normalized code $\hat{z}$:

$$\begin{aligned} F &= C(X), & \bar{z} &= \mathcal{N}(E(S)), \\ \hat{z} &= T(F, P), & \mathcal{L}_{\text{code}} &= \text{MSE}(\hat{z}, \bar{z}). \end{aligned} \quad (2)$$

$\bar{z}$ is the frozen teacher target; $\mathcal{N}$ is fixed channel normalization estimated from training-sheet AE codes. MSE averages over all code elements. Intermediate code heads receive the same target. A same-sheet loss pulls predictions from different prompt sets toward their mean; a different-sheet MSE-margin loss separates codes within the same CT crop.

The co-trained binary decoder B predicts union probability $V = \text{sigmoid}(B(F))$. BCE and Dice supervise V

against the GT sheet union Y on annotated voxels. This stage updates C , T and B while the AE stays frozen. Its base objective needs no full-resolution sheet decoding: prompt-dependent prediction and supervision act on the compact latent grid.

2.3. Inference: cluster codes, then decode sheets

Compute F once and sample seeds from the binary branch’s foreground. Predict a code for each seed and cluster using full-code MSE. For each retained cluster, combine consistent multi-point code predictions and decode a sheet:

$$\hat{S} = \mathbf{1}[\text{sigmoid}(D(\mathcal{N}^{-1}(\hat{z}))) \geq \tau]. \quad (3)$$

$\hat{S}$ is the binary sheet mask, $\mathbf{1}$ the indicator function, τ a probability threshold, and $\mathcal{N}^{-1}$ reverses normalization. Unstable clusters are discarded. Duplicate handling, component cleanup and probability-based overlap resolution produce final instance IDs. Sharing CT features and clustering before decoding avoids reconstructing a mask for every seed (Fig. 3). Dimensions, loss weights, sampling and cleanup are detailed in Supplementary.

3. Experiments

Setup. We use Kaggle-only checkpoint **0058** for the main comparisons and evaluate all **106 released test cubes**. Training-case IDs are excluded, but labels are now public and some crops were previously inspected; this is not a fresh blind submission. External-data checkpoints, training




Figure 2. **Predicting the sheet code.** Aligned teacher and CT/prompt paths produce the normalized target $\bar{z}$ and prediction $\hat{z}$, with MSE directly between them. Shared CT features also supervise the binary decoder against the GT union. Blue modules train; purple dashed modules are frozen, white panels are data, green objects are codes/features, and gold denotes losses. Sheet decoding is an inference-only extension. The shown eight-point prediction uses the displayed CT crop.

Method	Dice	SD ₂	Topo	VOI	LB
Winner	0.6015	0.8799	0.3769	0.5616	0.6176
PLatS	0.4262	0.7494	0.2644	0.5396	0.5305

Table 1. Original source-ignore erasure on all 106 released test cubes. Equal-case means; automatic inference, no GT prompts. LB is the public formula reproduced locally.

Method	Dice	SD ₂	Topo	VOI	LB
Winner	0.5460	0.8078	0.5694	0.5719	0.6537
PLatS	0.4046	0.7120	0.5820	0.5390	0.6124

Table 2. Restoration diagnostic on the same 106 cases, using the full-size 5,000-voxel gate. All three terms use restored predictions and erased GT. This is not official LB.

details and all 1/2/4/8-point results are in Supplementary. We evaluate automatic union quality, ignore-erasure sensitivity, prompted geometry and identity, then downstream surfaces. The reported automatic benchmark fixes a separate foreground proposer across checkpoints; these saved predictions appear in Fig. 3. The co-trained proposal path and new scratch recipe have no accuracy results here.

3.1. Binary segmentation and topology

We first apply the original leaderboard protocol: erase source-ignore voxels from prediction and GT, then compute the public formula [7]:

$$J = 0.35 \text{SD}_2 + 0.30 \text{Topo} + 0.35 \text{VOI}. \quad (4)$$

J is the combined score, SD_2 Surface Dice at two voxels, Topo exact matched TopoScore [6], and VOI the public variation-of-information score. Table 1 compares the released winner with automatic PLatS, using no GT prompts and equal weight per cube.

The winner achieves the higher combined score (0.6176 versus 0.5305) and higher Dice. Its binary mask can still

connect distinct sheets (Fig. 4, top). PLatS separates layers in this example while missing bands. Thus the example illustrates a topology failure mode, while the full benchmark identifies coverage and geometric accuracy as limitations.

3.2. Separating ignore cuts from model defects


Erasing ignored labels can create holes in an otherwise continuous prediction (Fig. 4, bottom). We therefore also report an *ignore-restoration diagnostic*: restore cut instances from their own pre-erasure predictions, exclude entirely ignored instances, and remove instances smaller than 5,000 voxels after restoration. The winner’s identities are raw connected components; PLatS provides sheet IDs. GT remains erased. Table 2 uses the same three formula terms. This is *not an official leaderboard score*: restored unknown foreground can still incur penalties. The exact rule and filtering controls are in Supplementary.

The winner still has the higher combined score (0.6537 versus 0.6124). PLatS has the higher restored topology term (0.5820 versus 0.5694), but lower surface and VOI scores. Restoration therefore diagnoses erasure sensitivity; it does


A • Seeds sampled throughout the 3D crop

Filled markers: $z=160$. Hollow markers: $z=148-172$ slab.

Foreground seeds




Full-code MSE




230 kept / 282 discarded

Clustered seeds




Decoded sheets




B • All seeds sampled on the displayed $z=160$ plane

Foreground seeds




Full-code MSE




238 kept / 274 discarded


Clustered seeds



Decoded sheets



Encode CT once
One code per seed



Groups → multi-point
prompts → code prediction

Frozen sheet decoder
+ instance cleanup

Same CT plane in both rows; 512 seeds per run. Instance colors are local to each row. Gray: discarded.

Figure 3. **Automatic instance discovery.** A: seeds sampled throughout the 3D crop, with nearby-slice markers displayed. B: all 512 seeds sampled on the exact displayed plane. Both runs use the same CT, foreground proposer and clustering/decoding settings; encoding and sheet reconstruction remain 3D. Full-code MSE matrices share a logarithmic scale. White outlines mark retained clusters; cyan dividers isolate discarded seeds. Colors link seed groups to decoded sheets within each row; gray marks discards.

not establish a new challenge ranking.

Topology before ignore processing. Table 3 counts tunnels (b_1) and cavities (b_2) in the full saved prediction unions, including unannotated regions. PLatS has lower mean and median tunnel counts, but its union contains cavities, whereas the winner’s contains none. These counts characterize prediction topology; they are not errors relative to GT or spatially matched TopoScores.

3.3. Sheet geometry and latent identity

To separate representation quality from automatic seed selection, we use nested GT prompts on 764 eligible annotated bands. Before cleanup or ignore erasure, 87.3% of eight-point predictions have one component and no tunnels or cavities. This is a topology statistic, not a geometric correctness rate. Table 4 measures alignment: increasing

Method	b_1 : tunnels		b_2 : cavities	
	Mean	Median	Mean	Median
Winner	40.302	20.5	0.000	0.0
PLatS	34.849	2.5	7.972	0.0

Table 3. Raw binary prediction unions on the same 106 native crops as Tables 1–2. Whole-volume Betti counts use V construction (six-connected foreground), with equal weight per case. No ignore erasure, restoration or additional size filtering is applied; each method’s existing inference postprocessing is retained.

prompts from one to eight raises Surface Dice from 0.7689 to 0.8058, while voxel Dice remains modest. GT components and repaired training targets favor simple sheet topology.

For 6,112 single-point codes, same-band versus




Figure 4. **Two different topology failures.** Top: a winner component joins six annotated bands in case 00860 ($z = 160$); PLatS preserves gaps but misses bands. Colors are IDs within each method. Bottom: ignore erasure cuts a predicted sheet in case 00829 ($z = 95$). Its whole-volume topology changes from one component and no tunnels to 495 components and 1,023 tunnels. Restoring its own voxels reverses the cuts. Orange marks removed voxels. Both rows use native z slices; selection rules and the enlarged orange-square crop are in Supplementary.

Points	Dice	SD ₂	Topo	VOI	Formula
1	0.4314	0.7689	0.6864	0.7502	0.7376
8	0.4517	0.8058	0.7045	0.7440	0.7538

Table 4. Prompted 0058 sheets with no cleanup. Only the exterior of the annotated-region box is ignored; internal unannotated voxels are background. All three formula terms are shown. This policy and GT prompts differ from the automatic original-erasure comparison in Table 1.

different-band classification has 0.9121 AUROC. The distributions overlap: at the inherited MSE threshold 0.02, 50.03% of same-band pairs exceed it. Automatic instance accuracy, $TP/(TP+FP+FN)$, is 0.420 at IoU 0.25 but only 0.071 at IoU 0.50. Useful code separation therefore coexists with clustering and alignment errors. Supplementary reports complete prompt counts, instance diagnostics and raw topology; none is a causal ablation of the AE or identity losses.

3.4. From sheets to ink

Predictions are meshed, independently flattened, and passed to a frozen ink model [1, 4, 5]. Figure 5 shows two PHERC0139 regions of approximately 20×33 and 23×44 mm. Each begins with one published-track point and grows using predicted points; reference layouts and ink labels are excluded from reconstruction. Median accepted join distances are 1.181 and 1.361 voxels, versus reference-distance p90 values of 15.856 and 15.234 voxels: local continuity does not ensure global alignment. Page A covers 17.37% of annotated pixels. Gaps and diffuse ink remain. These outputs demonstrate the pipeline, not new text recovery; the ink model’s training includes this scroll.

4. Discussion and conclusion

PLatS uses a shared sheet code for prompted reconstruction and automatic instance discovery. Raw topology and latent separability support this design, while union scores, strict instance matching and large-page drift expose its cur-




Figure 5. **Centimeter-scale surfaces and ink.** Independently flattened CT texture and ink probabilities from PLaTS on two PHerc0139 regions. Darker ink pixels indicate higher probability. Whole charts expose gaps and stitching limitations; native PNGs accompany the report.

rent limits. Matched direct-decoder and loss ablations are still needed to isolate the contribution of the bottleneck and identity supervision. Measured GPU time and memory are needed to quantify the benefit of clustering before decoding. Spatially disjoint scroll evaluation and broader delivered ink coverage are needed to establish generalization and reading benefit.

References

- [1] Hob3rMallow. ScrollFiesta: Virtual meshing and unwrapping for the herculaneum papyri. https://github.com/Hob3rMallow/scrollfiesta_public. Local extraction driver built from revision 5d957e9c21529580cc43eda8bf980fef02486b36.
- [2] Tony Li, Ozan M., Yiheng Wang, and Paul G. 1st place solution for the vesuvius challenge – surface detection competition. Kaggle solution writeup, February 27, 2026, 2026.
- [3] Ozan Oktay, Enzo Ferrante, Konstantinos Kamnitsas, Matthias Heinrich, Wenjia Bai, Jose Caballero, Stuart Cook, Antonio de Marvao, Timothy Dawes, Declan O'Regan, Bernhard Kainz, Ben Glocker, and Daniel Rueckert. Anatomically constrained neural networks (ACNN): Application to cardiac image enhancement and segmentation. *arXiv preprint arXiv:1705.08302*, 2017.
- [4] Michael Rabinovich, Roi Poranne, Daniele Panozzo, and Olga Sorkine-Hornung. Scalable locally injective mappings. *ACM Transactions on Graphics*, 2017.
- [5] ScrollPrize. 9- μ m ink detection model. https://huggingface.co/scrollprize/ink_9um. Hybrid 3D/2D checkpoint seed42, step 75000, revision 7109667e2607db1b90c37c8b09cb876ea7fe7bb1.
- [6] Nico Stucki, Johannes C. Paetzold, Suprosanna Shit, Björn Menze, and Ulrich Bauer. Topologically faithful image segmentation via induced matching of persistence barcodes. *arXiv preprint arXiv:2211.15272*, 2022.
- [7] Vesuvius Challenge. Surface detection: Evaluation. Kaggle competition and public metric resource package, 2026.
- [8] Li Yi, Wang Zhao, He Wang, Minhyuk Sung, and Leonidas Guibas. GSPN: Generative shape proposal network for 3D instance segmentation in point cloud. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2019.
- [9] Yuchen Zhou, Jiayuan Gu, Tung Yen Chiang, Fanbo Xiang, and Hao Su. Point-SAM: Promptable 3D segmentation model for point clouds. *arXiv preprint arXiv:2406.17741*, 2024.

Supplementary Material

A. Exact objectives and implementation

A.1. Autoencoder and full training objectives

Let X be a CT crop of $H \times W \times N_z$ voxels, S_i its binary mask for sheet i , and $P_i = \{p_k\}_{k=1}^K$ its K positive prompt coordinates. Here H, W, N_z are the three spatial dimensions. The AE encoder E maps a corrupted mask $\tilde{S}_i$ to a code z_i ; its occupancy decoder D predicts probability Q_i :

$$z_i = E(\tilde{S}_i), \quad Q_i = \text{sigmoid}(D(z_i \odot (1 + \xi))). \quad (5)$$

Here sigmoid is the logistic function, $\odot$ is elementwise multiplication, and ξ is independent uniform noise in $[-0.02, 0.02]$ during AE training. At 320^3 input, the code is 64×10^3 . No encoder–decoder skip connections bypass it (Fig. 1). The other heads predict a dense unsigned distance field U_i , a distance $h(z_i, q)$ at query coordinate q , and four intermediate occupancy maps $Q_i^{(a)}$.

Define $\mathcal{L}_{\text{seg}}$ as weighted BCE plus soft Dice loss. The AE uses

$$\begin{aligned} \mathcal{L}_{\text{AE}} = & \mathcal{L}_{\text{seg}}(Q_i, S_i) + \sum_a \alpha_a \mathcal{L}_{\text{seg}}(Q_i^{(a)}, S_i^{(a)}) \\ & + \lambda_U \mathcal{H}(U_i, d_i) + \lambda_q \mathbb{E}_q \mathcal{H}(h(z_i, q), d_i(q)) \\ & + \lambda_r \mathbb{E}_{i \neq j} [\text{ReLU}(\cos(z_i, z_j) - m)] \\ & + \lambda_o R_o(Q_i, S_i) + \lambda_b R_b(Q_i, S_i). \end{aligned} \quad (6)$$

The index a denotes decoder stages; $S_i^{(a)}$ is the resized target. $\mathcal{H}$ is Smooth L1 loss, d_i is unsigned distance to the sheet, clipped at 16 voxels and divided by 16, and $d_i(q)$ is its sampled value. The cosine term repels codes of distinct sheets in the same crop; $m = 0.5$ is its margin and $\text{ReLU}(t) = \max(t, 0)$. R_o and R_b penalize foreground outside a two-voxel GT tolerance and at the volume border. The α_a and λ symbols are loss weights, specified below. The sheet prior uses no explicit topology loss.

A.2. Point-conditioned model and joint binary head

The CT encoder C produces a 512×10^3 grid and four context attention blocks. Each prompt token combines Fourier coordinates, trilinearly sampled CT features and its positive label. Prefix attention and four transformer blocks T predict the normalized sheet code (Fig. 2):

$$\begin{aligned} \bar{z}_i &= (E(S_i) - \mu)/\sigma, \quad \hat{z}_i = T(C(X), P_i), \\ \hat{S}_i &= \mathbf{1}[\text{sigmoid}(D(\mu + \sigma \odot \hat{z}_i)) \geq 0.5]. \end{aligned} \quad (7)$$

μ and σ are fixed channel means and safe standard deviations of AE codes; $\mathbf{1}$ is an indicator. The AE remains frozen. A separate binary head B , with its own context refiner and convolutional upsampling, predicts sheet union probability

$V = \text{sigmoid}(B(C(X)))$ and supervises the shared CT features.

For transformer round r , let $a_i^{(r)}$ denote its auxiliary code, and let $d(u, v) = \|u - v\|_2^2 / L$ be MSE over L code elements. Training minimizes

$$\begin{aligned} \mathcal{L}_{\text{train}} = & d(\hat{z}_i, \bar{z}_i) + \sum_{r=1}^3 \beta_r d(a_i^{(r)}, \bar{z}_i) \\ & + 0.5 \mathcal{L}_{\text{same}} + 0.25 \mathcal{L}_{\text{diff}} \\ & + \mathcal{L}_{\text{seg}}^\Omega(V, Y), \end{aligned} \quad (8)$$

$$\begin{aligned} \mathcal{L}_{\text{same}} &= \mathbb{E}_{i,P} d(\hat{z}_i, P, c_i), \\ \mathcal{L}_{\text{diff}} &= \mathbb{E}_{i \neq j} [\text{ReLU}(1 - d(\hat{z}_i, \hat{z}_j))]. \end{aligned}$$

c_i is the mean code of same-sheet prompt sets P in a batch. Different-sheet pairs share a CT crop. $\beta = (0, 0.5, 1)$ weights the first three auxiliary heads. The final code uses a separate LayerNorm and linear head with weight 1; the instantiated fourth auxiliary head is unused by this loss configuration. Y is the GT binary union and Ω its non-ignored voxels. The binary loss uses positive BCE weight 2 and soft Dice on 160^3 crops. Point-to-sheet training uses no decoded occupancy loss in the main model.

A.3. Connections between the AE modules

Figure 6 expands the AE heads and their supervision. The reference dense encoder begins with a factor-four patch stem, followed by stage-matched residual convolutions and downsampling to a 10^3 grid. The code has 64 channels. The implementation also retains an occupied-voxel sparse encoder for future efficiency experiments; it densifies its bottleneck for the same decoder and heads. It is not used by the reported checkpoints.


The dense decoder applies five residual/convolutional upsampling stages, with spatial sizes 10, 20, 40, 80, 160, 320 and channels 64, 48, 32, 24, 16, 8. Final $1 \times 1 \times 1$ heads predict occupancy logits and unsigned distance. Four intermediate occupancy heads attach at spatial sizes 20, 40, 80, 160. The query head instead samples the noisy code at a query coordinate, concatenates normalized coordinates, and applies a width-128 MLP. There are no encoder–decoder skip connections.

A.4. Connections between the P2SD modules

Figure 7 separates shared image computation from prompt-dependent computation. The image encoder reduces CT to a 512×10^3 feature grid. A $1 \times 1 \times 1$ projection and four self-attention blocks form the image context F : 1,000 tokens of width 512. Three-dimensional per-axis rotary coordinates supply position information.

For each prompt coordinate p_k , prompt composition concatenates its eight-band Fourier encoding, the trilinearly sampled context $F(p_k)$, and the positive-label scalar. A

AE: learn a single-sheet code and freeze it for P2SD



Code regularization: distinct sheets in one crop $\rightarrow$ cosine repulsion on z (weight .1, margin .5).

All reconstruction heads depend on the bottleneck. No encoder-decoder skips; no explicit topology loss.

Figure 6. **AE modules and supervision.** Decoder occupancy and distance heads share dense features; query distance is predicted directly from the noisy latent code. Each loss is paired with its target. Intermediate heads attach at successive decoder stages. Code repulsion acts on distinct-sheet codes before latent noise. The sparse encoder remains an optional alternative.

two-layer MLP produces one width-512 token per point. The reference retains all K prompt tokens. The prompt modulator prefixes them to the context tokens, aligns their coordinate frames for rotary attention, and applies one attention block. Prompt-prefix outputs are discarded; the remaining 1,000 conditioned grid tokens enter the refiner.

The refiner contains four self-attention blocks with eight heads and SwiGLU MLPs. Linear projections supply intermediate code estimates; a separate LayerNorm and linear projection supply the final normalized 64-channel code. The loss distinguishes these auxiliary projections from the final head, as specified in Eq. 8. Channel statistics come from training-sheet AE codes. Inference reverses this normalization before the frozen AE decoder.

The binary branch receives the *unprompted* image context, applies its own four attention blocks, and up-samples with residual/convolutional stages of channels 512, 256, 128, 64, 32, 16. A final $1 \times 1 \times 1$ head produces union logits. It has no encoder-decoder skip connections; residual connections inside individual blocks are distinct. Its loss updates the shared image encoder and image-context blocks. In the new training recipe, refinement sees the complete 10^3 context and decoding builds the full 160^3 intermediate feature grid. We crop 80^3 features there and apply the final $\times 2$ upsampling, residual block and occupancy head to obtain 160^3 output supervision. Earlier stages retain full-volume context at increased activation-memory

cost. This is the retained `second_last` crop mode, also used by 0076. The scored 0058 checkpoint instead used an earlier 5^3 context crop; its measurements are unchanged.

A.5. Efficiency of compact latent prediction

For a 320^3 crop, prompt conditioning and code prediction operate on only $10^3 = 1,000$ spatial sites, compared with 32,768,000 output voxels. Including its 64 channels, a sheet code has 64,000 scalars, 512 times fewer than one full-resolution scalar mask. At BF16, these tensors occupy 0.122 MiB and 62.5 MiB, respectively; the shared 512×10^3 context occupies 0.977 MiB. These are tensor-size comparisons, not measured speedups or total training-memory estimates. Attention still incurs token-pair work.

Let N be the number of foreground seeds, G the number of additional multi-point prompt groups used for cluster consistency, and K the number of sheet hypotheses that are densely decoded. Let C_{img} denote shared CT encoding/context cost, C_{code} the cost per prompt group, C_{cluster} clustering cost, C_B the binary branch’s refinement and union decoding cost, and C_D one AE decoding cost. With the co-trained proposal head, automatic inference has approximate cost

$$C_{\text{img}} + C_B + (N + G)C_{\text{code}} + C_{\text{cluster}} + KC_D. \quad (9)$$

This avoids paying a full upsampling cost for each of the N seed codes before clustering. Retained sheets still require

P2SD: new training recipe with late-stage binary cropping




Figure 7. **P2SD modules in the new training recipe.** Shared CT encoding and context feed prompt composition, modulation and refinement. Code heads learn from the frozen AE. The binary branch decodes full context to 160³ features before cropping for its final upsampling; at inference it decodes the full union for seed proposal. The sheet decoder is used only when masks are needed under the main latent objective. Optional query losses are described separately; the scored 0058 checkpoint used earlier cropping.

dense decoding, output storage and cleanup. The co-trained proposal head reuses the cached image context and incurs one full-resolution union decode. The historical benchmark instead uses a separate proposer with its own CT encoder, so its cost includes an additional encoding. No binary training crop is applied at inference.

During P2SD training, the frozen teacher encoder produces sheet targets, but the active latent objective does not require AE decoding. The auxiliary binary branch still decodes a crop once per CT batch, and validation decodes full sheet masks. We therefore claim an architectural advantage in repeated prompt processing, rather than an unmeasured end-to-end runtime advantage.

A.6. Retained query supervision

An optional P2SD-owned distance head h_θ supplies geometric supervision without full sheet decoding. It concatenates a trilinear sample of the raw predicted code $\hat{z}_i = \mu + \sigma \odot \hat{z}_i$ with normalized query coordinates q , and ap-

plies a width-128 MLP. Its direct GT loss is

$$\mathcal{L}_{\text{query}}^{\text{GT}} = \mathbb{E}_q \mathcal{H}(h_\theta(\hat{z}_i, q), d_i(q)). \quad (10)$$

Here θ denotes the trainable query-head parameters, $\mathcal{H}$ is Smooth L1, and $d_i(q)$ is the clipped, normalized GT distance defined above. The optional term updates both the head and the predicted code. Direct GT targets avoid inheriting errors from the AE query head. A separate retained distillation option uses frozen-AE query predictions as targets instead.

The P2SD-owned head currently predicts distance only. Query occupancy/distance heads and GT-target utilities remain available in the implicit AE path; a separate GT-supervised P2SD occupancy-query head is not wired in this release. These optional paths are preserved for experiments, not included in the 0058 objective or enabled by the new default scratch recipe.

A.7. AE flip/rotation consistency: proposed extension

The same geometric idea used in the P2SD continuation can be applied to the AE, but this extension is *not implemented or evaluated here*. Define $\mathcal{A} = \text{sigmoid} \circ D \circ E$ with latent noise disabled during the consistency passes. For an exact flip or right-angle rotation g , a candidate output-space objective is

$$\mathcal{L}_{\text{AE-eq}} = \text{MSE}(\mathcal{A}(g\tilde{S}), g\mathcal{A}(\tilde{S})). \quad (11)$$

$\tilde{S}$ is one corrupted sheet; g transforms both the input and output voxel grids, and MSE averages over output voxels. Both branches use the same transformed corruption and retain supervision by the corresponding clean GT masks. A small weight after reconstruction stabilizes, applied periodically, would limit the additional decoder cost. This does not assume equivariance of the latent channels. A retrained AE would require fresh latent statistics and P2SD training against its updated codes.

A.8. Loss definitions and architecture settings

For a probability map Q , binary target S , valid voxel set Ω and positive BCE weight ω , the segmentation loss in the main text is

$$\begin{aligned} \mathcal{L}_{\text{seg}}^\Omega &= \mathcal{L}_{\text{BCE}}^\Omega + \mathcal{L}_{\text{Dice}}^\Omega, \\ \mathcal{L}_{\text{BCE}}^\Omega &= -\frac{1}{|\Omega|} \sum_{v \in \Omega} [\omega S_v \log Q_v + (1 - S_v) \log(1 - Q_v)], \\ \mathcal{L}_{\text{Dice}}^\Omega &= 1 - \frac{2 \sum_{v \in \Omega} Q_v S_v + \epsilon}{\sum_{v \in \Omega} Q_v + \sum_{v \in \Omega} S_v + \epsilon}. \end{aligned} \quad (12)$$

$Q_v = Q(v)$ and $S_v = S(v)$; v indexes voxels and $\epsilon = 10^{-6}$ stabilizes the Dice ratio. Without the superscript, all voxels are valid. AE BCE uses $\omega = \min(20, \max(1, N_- / \max(N_+, 1)))$, where N_+ and N_- count foreground/background voxels in the training batch. Intermediate targets use max pooling to the head resolution. The joint binary head uses $\omega = 2$; ignore is removed from both Dice sums, and empty-target crops contribute BCE but are excluded from the batch Dice average.

In Eq. 6, the four intermediate weights are $\alpha = (0.125, 0.25, 0.5, 0.5)$ and $(\lambda_U, \lambda_q, \lambda_r, \lambda_o, \lambda_b) = (0.15, 0.15, 0.1, 0.05, 0.05)$. Smooth L1 uses transition width 0.1 and averages over dense voxels or query samples. The penalties are

$$R_o = \frac{\sum_{v \notin \text{dilate}_2(S)} Q(v)}{\max(|S|, 1)}, \quad R_b = \frac{\sum_{v \in \partial_5} Q(v)}{\max(|S|, 1)}. \quad (13)$$

dilate_2 is a cubic two-voxel dilation and ∂_5 is the five-voxel crop border. Cosine similarity flattens the latent tensor. AE input corruption is applied with probability 0.7: one to six

spherical holes of radius 2–10 voxels and dropout up to 0.3, retaining at least 25% of the sheet. The unsigned-distance query head uses 8,192 queries.

The AE uses convolutional channels (8, 16, 24, 32, 48, 64), RMS normalization, and factor-32 reduction; its distance-query MLP has width 128. The CT encoder channels are (16, 32, 96, 192, 384, 512), with an early factor-four stem. Point Fourier encoding has eight bands. Context and point-refinement attention each have four blocks, 512-dimensional features and per-axis rotary coordinates; refinement uses eight heads. The joint binary branch has its own four-block refiner and convolutional decoder channels (512, 256, 128, 64, 32, 16). Historical 0058 binary training decodes 5^3 context crops into 160^3 masks; 0076 and the new scratch recipe use the late feature crop in Fig. 7. Prompted sheet codes retain the full 320^3 crop extent. Checkpoint, configuration and latent-statistics hashes accompany the result snapshot. The AE and main CT model use Kaggle-derived t6 training labels with topology repair/cleanup; external scroll labels are excluded from the main model. Evaluation uses the released test annotations, not those repaired training labels. Training-ID exclusion alone does not certify spatial disjointness. Cached CT inputs share five-voxel image-border zeroing; GT borders are retained.

B. Automatic clustering and instance diagnostics

Figures 1–2: real model examples. All image panels use case 00860 at native $z = 160$. The AE input, output and target show the identical 160×160 crop ($x = 80:240$, $y = 40:200$) of annotated band 9. Corruption follows the original AE training recipe with torch seed 42; occupancy is displayed at threshold 0.5 without cleanup. Dense distance panels share the training scale: distance clipped at 16 voxels and divided by 16. Coordinate-query supervision samples that same 3D target. The two-sheet example encodes clean bands 9 and 8; their cosine is 0.332. This is an illustration, not a separation benchmark or causal ablation. Code glyphs denote whole tensors; they are not feature-intensity measurements.

The prompted path uses the frozen eight-point manifest for band 9 and checkpoint 0058. CT and output show the full native plane; only prompts within $z \pm 12$ are overlaid as hollow off-plane markers (their exact z values are 153 and 151). All eight 3D prompts are used for prediction. The depicted binary decoder contains residual blocks but has no encoder–decoder skip connections. Intermediate code heads and identity losses are defined above; the overview emphasizes the target/prediction MSE and the separate inference-only decode.

Figure 3: measured automatic clustering. Row A uses the original 512 automatic seeds and final instances. The saved codes retain all 64,000 dimensions. Reapplying the original leader clustering threshold reproduces every final seed assignment and all 18 retained cluster sizes. Of 112 initial clusters, 94 contain at most eight points and are excluded: 230 seeds remain and 282 are discarded. Duplicate handling maps the 18 retained clusters to 14 final IDs. Matrix rows and columns are ordered by final ID and then leader-cluster ID, with discarded clusters last. Colored strips denote final IDs; gray denotes discarded points. White outlines bound retained clusters, and cyan lines divide retained from discarded seeds.

Distances are full-code MSE, computed in double precision from the saved codes; the logarithmic display floors values below 10^{-5} only for coloring. The matrix visualizes relationships, while assignment uses distance to a running cluster centroid. CT seed markers use the explicitly labeled $z = 148:172$ slab: filled markers lie on $z = 160$, hollow markers lie off-plane. The CT and decoded-mask panels use the exact $z = 160$ section. Seed colors are assigned after inference for display, never supplied to the model. Within each retained cluster, multi-point prompt subsets use the cached CT context to predict codes; consistent subset codes are combined before decoding and cleanup.

Exact-plane seed example (Figure 3B). We restrict seed support to the cached foreground on the same native $z = 160$ plane and sample 512 distinct points uniformly without replacement (seed 0). Every seed and every multi-point prompt lies on the displayed plane. The CT input, validity extent, code prediction and sheet decoding remain fully 3D; only seed support changes. The same checkpoint and automatic settings produce 93 initial clusters, 17 retained clusters and 13 final instances. After exclusions, 238 seeds remain and 274 are discarded. No GT or ignore masks are loaded. This illustrative run is separate from the fixed benchmark tables. Both matrix panels use the same logarithmic color scale; instance colors are local to each row and are not a cross-run identity matching.

Measured latent identity separation. Each of the eight fixed GT coordinates per band is queried separately, yielding 6,112 single-point codes across all 106 cubes. Distance uses the complete normalized code, without decoding or threshold calibration. GT connected bands provide analysis identities, rather than physical-sheet guarantees. On 106 cases and 6,112 fixed GT points, the median same-band MSE is 0.0201, versus 2.1280 between bands. Pair AUROC is 0.9121; mean per-case nearest-band accuracy is 80.80%. At the inherited 0.02 threshold, 50.03% of same-band pairs exceed the threshold and 1.31% of different-band pairs fall below it. The same-band p90 is 1.682, so the tails

overlap substantially despite separated medians. These spatially spread GT points include difficult crop-edge locations and differ from the automatic foreground-seed distribution. Pairwise rejection is not the fraction of points excluded by leader clustering; nearest-neighbor links and cluster means affect assignment. Results support useful separation while identifying remaining point-conditioning errors.

The controlled foreground proposer is frozen model 0022, shared across all checkpoints; it is separate from their jointly trained binary heads. Foreground probability threshold is 0.6, with 512 uniformly sampled voxels. One cached CT context serves all single-point latent predictions. Sequential leader clustering joins the nearest cluster mean at latent MSE ≤ 0.02 ; otherwise it creates a new cluster. Clusters with eight or fewer points are discarded. For seven eight-point subset predictions u_t , define

$$r_t = \min_{s \neq t} d(u_t, u_s), \quad r_{\min} = \min_t r_t. \quad (14)$$

Here $s, t \in \{1, \dots, 7\}$ index subsets, and d is normalized latent MSE as in the main text. Candidates with $r_{\min} > 0.008$ are discarded; otherwise all seven codes are averaged and decoded once at threshold 0.5. The rule checks subset agreement, not certified physical identity.

Each automatic decode retains its largest six-connected component. Three-voxel dilation with overlap-cover threshold 0.2 suppresses duplicate masks; predicted probability assigns contested voxels, and final IDs below 500 voxels are removed. Thus the automatic benchmark includes cleanup, whereas raw prompted topology does not. Both small and unstable point clusters are excluded from reconstruction. Leader order and fixed latent thresholds remain limitations.

Tables 5 and 6 retain all automatic checkpoint results on the 106 cases. The unfiltered common-foreground CC6 baseline is not a matched-cleanup architecture ablation. Instance accuracy at IoU threshold τ is $TP/(TP+FP+FN)$, using greedy one-to-one annotated-band matching; TP, FP and FN are matched predictions, unmatched predictions and missed bands. It is not integrated AP. GT-normalized IoU sums matches at 0.5 and assigns zero to unmatched GT bands. Merges count predicted majority-overlap claimants for multiple GT bands; splits count GT bands receiving multiple majority claimants. Annotation components are proxies, not certified physical-sheet identities.

C. External data and continued training

The two additional checkpoints investigate whether large-scale external data and prolonged training help generalization to new scrolls. **0076** initializes from Kaggle-only 0058 and adds mesh-derived scroll patches, totaling 10,347 spatially overlapping training patches. It adds latent coordinate-distance supervision, with query emphasis near mesh vertices. **0076+100k** continues for 100,000 steps

Arm	Dice	Surface Dice	TopoScore	VOI	Formula
Foreground CC6	0.4191	0.7377	0.0883	0.5358	0.4722
0058	0.4262	0.7494	0.2644	0.5396	0.5305
0076	0.3847	0.6863	0.2269	0.5426	0.4982
0076+100k	0.3981	0.7098	0.2386	0.5403	0.5091

Table 5. Completed automatic segmentation on 106 cubes under original source ignore. One common foreground and the same 512 sampled points are used for every P2SD checkpoint. The CC6 baseline labels the foreground directly.

Arm	Acc. _{.25}	Acc. _{.50}	GT IoU. _{.50}	Pred	GT	Merges	Splits
Foreground CC6	0.031	0.007	0.038	299.953	7.642	1.189	0.000
0058	0.420	0.071	0.057	6.821	7.642	0.330	0.000
0076	0.386	0.064	0.048	6.217	7.642	0.302	0.000
0076+100k	0.385	0.059	0.046	6.396	7.642	0.321	0.000

Table 6. Per-case means of annotated-band instance diagnostics. Acc_τ is $\text{TP}/(\text{TP}+\text{FP}+\text{FN})$ at IoU τ , not integrated AP. GT IoU sums IoUs of matches at threshold 0.50 and divides by the GT count, assigning zero to unmatched bands. Pred/GT are counts; merge and split definitions are those of the inherited matching implementation.

with geometric consistency. All three retain the same frozen AE. Changed data, objectives and training duration prevent a causal attribution to any one factor.

Both mixed-data checkpoints enable the P2SD-owned GT-distance term in Eq. 10, with weight 0.15, 4,096 sampled queries and additional mesh-vertex distance targets. Their frozen-AE query-distillation option is disabled. The continuation additionally compares aligned decoded probabilities under flips and right-angle rotations:

$$\mathcal{L}_{\text{geo}} = 1 - \frac{2\langle A_g, B_g \rangle + 1}{\|A_g\|_2^2 + \|B_g\|_2^2 + 1}. \quad (15)$$

A_g and B_g are original/transformed 160^3 probability blocks after inverse alignment; inner products and norms sum voxels. A nonidentity transform is sampled each step from the 16-element in-plane dihedral and axial-flip group. Corresponding 5^3 code crops are decoded, with gradients through both branches and the frozen decoder. The consistency weight is 0.5. AdamW uses stem, encoder, dense-head and other learning rates 10^{-6} , 3×10^{-6} , 10^{-5} , 5×10^{-6} , weight decay 0.02, 1,000 warmup steps and cosine decay to 10% of their initial rates. These settings allow intermediate checkpoint evaluation without relying on a near-zero terminal learning rate.

D. Full prompt-count and topology results

Of the 106 released test cubes, 102 are 320^3 and four are 256^3 . Training-case IDs are disjoint; this does not establish spatial disjointness within a scroll. The first-place baseline uses its released four-model ensemble and postprocessing unchanged; only GPU routing is adapted. Across 764 raw eight-point sheet predictions, mean (b_0, b_1, b_2) is

$(1.183, 0.165, 0)$. After source-ignore erasure, mean b_0 rises to 22.857 and mean b_1 to 32.191. These whole-mask counts differ from matched TopoScore. Here b_0, b_1, b_2 count components, tunnels and cavities under six-connected foreground. In the selected binary failure crop, the raw winner union has 125 tunnels and the automatic PLaS union has three; this example is diagnostic, not a representative estimate.

Eight spatially spread positive GT points are fixed per eligible annotated band, with nested prefixes $K \in \{1, 2, 4, 8\}$ identical across checkpoints. The complete matrix has $3 \times 4 \times 764 = 9,168$ sheet tuples and 1,272 checkpoint/prompt/case unions. Smaller native cubes are centered with 32-voxel padding; masks, points and scores are cropped back. Every tuple has aligned image, prompt, GT, prediction and ignore NIFTIs.

The prompted diagnostic ignores only the exterior of the rectangular annotated-region box, retaining source-ignore regions *inside* that box as GT background. This avoids inward erasure cuts but is distinct from the standard LB and restoration tables in the main text. Raw masks receive no cleanup. Union GT includes annotated bands below the 10,000-voxel prompt eligibility threshold. Tables 7 and 8 retain Dice and all three formula terms for every prompt count.

Table 9 counts raw topology before erasure. All three median per-sheet vectors are $(1, 0, 0)$, while mean union tunnel counts expose outlier contacts between decoded sheets. For 0058 the mean raw union b_1 is 63.292 but its median is one: mean counts must not be read as typical defects per crop. The GT component construction also favors disk-like masks. P_{100} denotes fraction $(1, 0, 0)$ and is not proof of geometric correctness. Exact TopoScore uses spa-

Checkpoint	Points	Dice	SD ₂	Topo	VOI	Formula
0058	1	0.4314	0.7689	0.6864	0.7502	0.7376
0058	2	0.4329	0.7735	0.6994	0.7507	0.7433
0058	4	0.4471	0.7936	0.6954	0.7463	0.7476
0058	8	0.4517	0.8058	0.7045	0.7440	0.7538
0076	1	0.4050	0.7445	0.6576	0.7557	0.7224
0076	2	0.4115	0.7564	0.6546	0.7530	0.7247
0076	4	0.4287	0.7820	0.6635	0.7477	0.7344
0076	8	0.4334	0.7987	0.6772	0.7450	0.7435
0076+100k	1	0.4214	0.7621	0.6729	0.7505	0.7313
0076+100k	2	0.4269	0.7718	0.6658	0.7485	0.7319
0076+100k	4	0.4413	0.7939	0.6719	0.7446	0.7400
0076+100k	8	0.4422	0.8071	0.6821	0.7430	0.7472

Table 7. Complete sheet results: all three checkpoints, all four nested prompt counts. Annotated-box ignore; no prediction cleanup. Formula scores are diagnostics under this policy.

Checkpoint	Points	Dice	SD ₂	Topo	VOI	Formula
0058	1	0.4437	0.7922	0.5848	0.5336	0.6395
0058	2	0.4465	0.7978	0.6054	0.5352	0.6482
0058	4	0.4582	0.8132	0.6264	0.5343	0.6595
0058	8	0.4619	0.8221	0.6171	0.5341	0.6598
0076	1	0.4166	0.7692	0.5597	0.5296	0.6225
0076	2	0.4220	0.7787	0.5373	0.5322	0.6200
0076	4	0.4368	0.7994	0.5704	0.5335	0.6376
0076	8	0.4426	0.8122	0.5990	0.5305	0.6497
0076+100k	1	0.4313	0.7836	0.5781	0.5308	0.6335
0076+100k	2	0.4370	0.7928	0.5543	0.5342	0.6308
0076+100k	4	0.4495	0.8114	0.5725	0.5349	0.6430
0076+100k	8	0.4512	0.8210	0.6065	0.5313	0.6553

Table 8. Complete union results: all three checkpoints, all four nested prompt counts. Annotated-box ignore; no prediction cleanup. Formula scores are diagnostics under this policy.

Checkpoint	Per-sheet means				Union means		
	b_0	b_1	b_2	P_{100}	b_0	b_1	b_2
0058	1.183	0.165	0.000	87.30%	8.179	63.292	9.528
0076	1.385	0.658	0.263	81.94%	9.594	19.906	3.557
0076+100k	1.308	0.534	0.211	82.46%	8.953	40.302	5.896
GT	1.000	0.001	0.001	99.74%	7.642	0.009	0.009

Table 9. Whole-mask raw Betti counts before scoring erasure or component cleanup. Sheet statistics cover 764 eligible bands; union statistics cover 106 cases. P_{100} denotes the fraction of sheets with $(b_0, b_1, b_2) = (1, 0, 0)$. These counts do not replace spatial feature matching.

tial Betti matching over eight crop octants, not a comparison of whole-mask counts alone. Our compact exact implementation reduces storage without approximating the matching calculation. Formula arithmetic, 106-case coverage, native prompt membership and unchanged checkpoint hashes are verified. Paired bootstrap results average sheets within each case and resample 106 cases 2,000 times (seed 42); source files preserve all intervals. Spatial dependence and prior in-

spection limit their statistical interpretation.

Raw automatic union census. Table 3 uses the winner’s saved binary mask and the union of automatic 0058 instance IDs, exactly as supplied to the original-erasure comparison. This differs from the GT-prompted sheet/union census above. We recompute all 212 native volumes without loading GT or ignore arrays, excluding no region and applying

no additional component filter. Existing inference cleanup remains part of each method’s prediction.

Under V construction, b_0 counts six-connected foreground components and b_2 counts enclosed 26-connected background components; outside the crop is background. With Euler characteristic χ , the tunnel count is $b_1 = b_0 + b_2 - \chi$. All counts agree with the previous per-case raw audit, and all prediction hashes match their inference freezes. PLaTS has nonzero b_2 in 31 of 106 cases (845 cavities in total); the winner has none. Means and medians expose this uneven distribution. The standalone census script is included with the report; per-case counts and input hashes are retained in the local experiment record.

E. Restoration controls and failure provenance

Preserve each predicted identity *before* erasure. Let A_i be its full mask, I the source-ignore mask, and $v_i = |A_i \setminus I|$ its visible size. The restored, filtered binary union is

$$A^{\text{rest}} = \bigcup_{i: v_i > 0, |A_i| \geq 5000} A_i, \quad J_{\text{rest}} = J(A^{\text{rest}}, G \setminus I). \quad (16)$$

$|\cdot|$ denotes voxel count. Entirely ignored instances are excluded; the 5,000-voxel gate uses *full size after restoration*. GT remains erased, and restored predictions are not erased a second time. For the binary winner, identities are its raw six-connected components; PLaTS supplies predicted IDs. The rule uses no GT foreground to select or complete a prediction. Table 2 reports the same three formula terms after this change. This is *not official LB*: restored foreground can be penalized against unknown GT background, and raw binary components are imperfect sheet proxies. Filtering-only controls and a topology-isolation diagnostic are in Supplementary.

The filtering-only control erases ignore after applying the same full-size ID gate. The full restoration metric is Eq. 16, including any Surface Dice or VOI penalties for retained unknown foreground. To isolate the changed topology term, we additionally report

$$\begin{aligned} J_{\text{topo-only}} = & 0.35 \text{SD}_2(A^{\text{rest}} \setminus I, G') \\ & + 0.30 \text{Topo}(A^{\text{rest}}, G') \\ & + 0.35 \text{VOI}(A^{\text{rest}} \setminus I, G'). \end{aligned} \quad (17)$$

where $G' = G \setminus I$. This hybrid is a separate non-official diagnostic; it is not the modified LB table in the main paper.

Failure examples use saved masks, never synthetic surfaces. The binary example uses a visually clear lamellar CT cross-section among the top 20 cached winner merger candidates. The central TIFF-stack slice is used, with a four-connected winner region joining distinct GT bands. The ignore example maximizes induced tunnel count among raw (1, 0, 0) 0058 masks with visible support. Source hashes,

selection metadata, native voxels and slice coordinates are saved in the failure-analysis run. Selection is diagnostic, not random or representative. Figures 4 and 8 show native 2D slices with nearest-neighbor display, using the same plane and crop across methods. TIFF arrays are $[z, y, x]$ and project arrays are $[x, y, z]$; a displayed z slice is the transpose of internal array $[:, :, z]$. CT panels are verified directly against the original TIFF pages. The ignore-cut slice maximizes disconnected visible components among continuous raw z sections with at least 1,000 foreground and 20 removed voxels. Counts in the main caption use the exact whole-volume voxel V construction; a displayed 2D interruption alone does not establish a 3D tunnel.

F. Large pages and spiral-fitting integration

The large-page main-paper figure uses 0058 on the same two preselected PHerc0139 regions as our earlier downstream control. Native sampling is 9.362 μm ; patch size 320 and step 128 give 60% linear overlap. The first decoded prediction defines a tangent frame; later points come from accepted predictions. An overlap acceptance gate can reject ambiguous growth, so accepted join error is conditional on acceptance and is not a global geometric error. Independent SLIM retains holes in predicted chart support. The inward normal is fixed from the scroll umbilicus before ink inference; full inward/outward PNGs are preserved. Geometry and partial ink labels enter only after all reconstruction/ink hashes are frozen. The known ink-model scroll and potentially overlapping external training regions are not held-out transcription tests. On page A, the single-method supported-pixel ink AP/Dice are 0.9170/0.8294; labeled coverage is 17.37% and delivered Dice is 0.2932. These supported pixels are not the common-pixel set used in the integration comparison below.

An earlier integration pilot, using original 0076, compares the direct prediction, published m7 tracks passed to the official spiral fitter, PLaTS patches alone, and tracks plus PLaTS patches. All fits have the same 30,000-step budget, canonical initialization, independent flattening and frozen ink. The published tracks have *unverified Kaggle-winner provenance*; this is not the first-place model comparison. Two predicted local charts alone leave global fitting underconstrained. We retain these controls to demonstrate an integration route, not an upper bound on the published fitting method.

Ink AP and Dice in Table 12 use the same 26,296 fully supported labeled pixels across all four arms. Labeled coverage and “delivered” Dice also count missing output, assigning zero to missing predictions. Pairwise comparisons use different common pixel sets and must not be mixed with the all-arm table. Page B has no audited ink labels. High common-pixel scores on a small delivered fraction do not establish whole-page letter recovery. This pilot studies in-

Method	Standard LB	Filter only	Restore + filter	Topology-only diagnostic
Winner	0.6176	0.6176	0.6537	0.6753
PLatS auto	0.5305	0.5304	0.6124	0.6257
PLatS GT8 union	0.5627	0.5627	0.6598	0.6611

Table 10. Ignore-restoration controls. GT8 union uses deterministic sheet-ID ownership on overlaps and retains its original union exactly. It uses GT prompts and is not an automatic discovery result. The topology-only hybrid replaces only TopoScore.




Figure 8. Enlarged native slice from the main-paper binary failure example: the same 90×90 -pixel crop at TIFF $z = 160$ in all four columns. The winner component connects distinct GT bands; automatic PLatS has lower coverage. Twenty-seven predicted voxels in the sampled gap are known background; none of its 34 samples are ignored. The red dashed segment marks this sampled gap in GT and winner panels. Labels are displayed without dilation or smoothing.

Page	Area cm^2	Accepted/attempted	Join median (vox)	Target median	Target p90
page_a	5.67	255/365	1.181	3.905	15.856
page_b	8.92	411/565	1.361	2.778	15.234

Table 11. 0058 large-page outputs. Areas are reconstructed mesh areas; joins include accepted windows only. Target distances exclude reference-boundary samples and are measured after all ink outputs are frozen.

Input/method	Area cm^2	Median vox	p90 vox	Ink AP	Ink Dice	Coverage	Delivered Dice
Direct PLatS	6.258	2.249	8.059	0.7582	0.6021	0.2348	0.3324
Published tracks	5.773	14.450	29.358	0.3314	0.1731	0.1598	0.0524
PLatS patches only	6.236	6.311	94.290	0.8393	0.7446	0.1444	0.1314
Tracks + PLatS	5.948	3.112	12.187	0.9881	0.9336	0.1591	0.2239

Table 12. Original 0076 integration pilot on page A. Common ink uses identical pixels across all four arms; coverage and delivered Dice expose missing output. Published tracks are not verified as the Kaggle winner.

tegration and drift.




Figure 9. Large-page spiral integration control on PHerc0139 page A, original 0076 (Supplementary only). Direct prediction, published-track fitting, prediction-only fitting and mixed-input fitting all use independent layouts and the same ink model. The track source is not verified as the Kaggle winner.